


## CS 425 Distributed Systems MP1 Report

Aryaman Ramchandran, Aditi Sadwelkar

Our distributed grep system implements a client-server architecture using socket programming adapted from Beej's guide. The client employs Python's ThreadPoolExecutor to query multiple VMs concurrently, sending JSON-formatted grep requests containing patterns and options. Each server executes grep locally on its designated log file (machine.X.log) and returns results.

Unit tests covered three query patterns across 4 VMs, each storing 60MB log files:

- (1) Broad case-insensitive search for "GET" representing high-match scenarios.
  - (2) Narrow date-specific search for "17/Aug" testing selective queries.
  - (3) Complex regex pattern "[0-9]{2}:[0-9]{2}:[0-9]{2}" evaluating computational overhead.
- Each test executed 10 trials to ensure statistical significance.



The performance analysis reveals distinct latency characteristics based on query complexity. The narrow search achieved the lowest average latency (149.9ms,  $\sigma=5.5$ ms) due to fewer matching lines reducing network transfer overhead.

The broad search showed moderate latency (416.3ms,  $\sigma=23.1$ ms) as expected from processing numerous matches.

The regex query exhibited the highest latency (624.4ms,  $\sigma=27.9$ ms), demonstrating the computational cost of pattern matching.

The low standard deviations (1.3-4.5% of means) indicate consistent system performance. The linear relationship between match volume and latency confirms that network I/O dominates execution time. Concurrent querying via ThreadPoolExecutor effectively masks individual server processing delays, achieving sub-second response times even for complex patterns across distributed logs.