

# CS425 Distributed Systems MP3 Report

Aditi Sadwelkar, Aryaman Ramchandran

## a) Design

### Replication Strategy

Files are mapped to ring positions using SHA-1 (secure hashing algorithm), and replicas are stored on the first 3 successor nodes on the ring to tolerate up to 2 simultaneous machine failures. The `ring.py` module implements the consistent hash ring where both nodes and files are hashed to positions. When a file is created or accessed, the system calls `get_successors` to determine the exact 3 nodes that should store the file. Node IDs follow the format `hostname:port:timestamp` to ensure uniqueness across node restarts.

### Consistency Guarantees Implementation

Each client maintains a sequence counter that increments with each append. When blocks are merged, sorting by `client_id`, `sequence_num`, `timestamp` ensures that appends from the same client always appear in the order they were issued.

The merge operation collects blocks from all replicas, deduplicates them using `block_id` as the unique identifier, sorts them to establish a total order, and redistributes the merged result to all replicas.

The ConsistencyManager tracks pending writes per `client_id`, `filename` pair. When a client issues a GET request, the system checks `has_pending_writes()` and blocks the read if that client has outstanding appends to the file, ensuring the client always sees its own writes.

### Re-Replication Implementation

The ReplicationManager runs a monitoring thread that checks every 5 seconds for under-replicated files. The `check_and_rereplicate()` method iterates through all stored files, queries each replica to determine actual replication count, and triggers re-replication if the count falls below 3. When a node failure is detected through MP2, the `handle_node_failure()` method is called, which removes the failed node from the ring and triggers the monitoring thread to fix under-replicated files.

### Merge Algorithm

The merge operation is coordinated by the first replica in the successor list. The coordinator sends requests to all replicas, including itself, collects all blocks, and passes them to `consistency.merge_blocks()`. This method deduplicates blocks, sorts, and validates ordering using `validate_block_ordering()` to ensure no sequence gaps exist per client. The merged blocks are then sent to all replicas via messages with the `replace=True` flag, causing each replica to delete old blocks and write the new merged set atomically.

# CS425 Distributed Systems MP3 Report

Aditi Sadwelkar, Aryaman Ramchandran

## b) Past MP Use

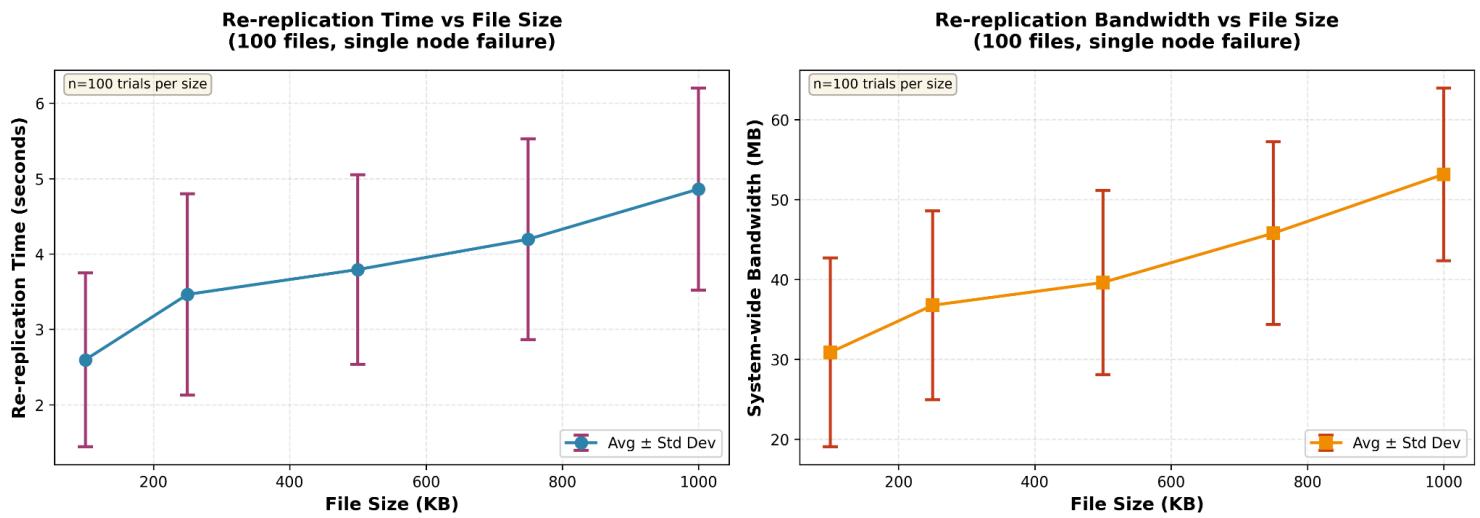
MP2's membership protocol is integrated directly into node.py through the MembershipService class and provides the foundation for our HyDFS's failure detection and node discovery. When nodes join or fail, MP2's PinkAck protocol detects the change, and HyDFS triggers rebalancing or re-replication.

MP1's distributed grep was useful for debugging by allowing queries like grep "REPLICATION ERROR" -i across all VM logs to quickly identify issues with file placement, merge conflicts, or network errors during development and testing.

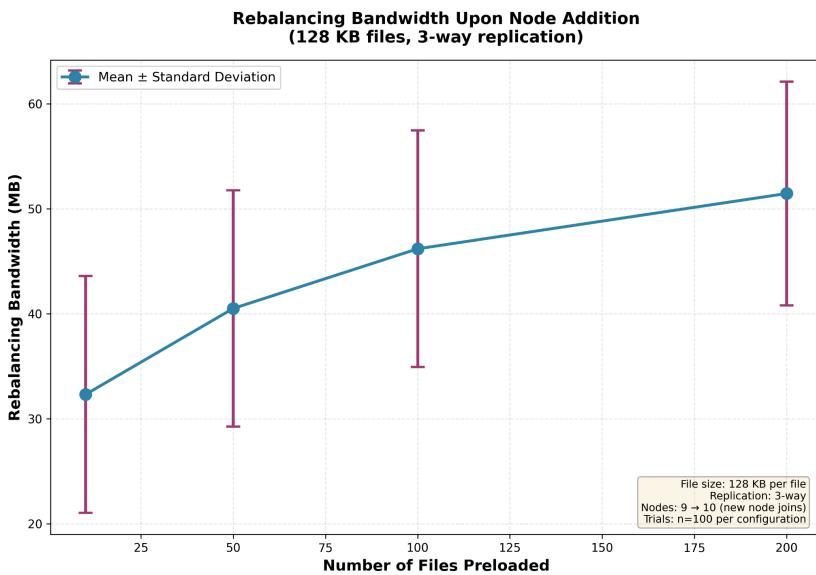
## c) LLM Use

LLM's have been used to bounce back and forth ideas regarding system design and implementation. Conceptual struggles and weighing strategy options to help make clear and informed decisions on direction. The LLM logs are stored in <llm\_logs> directory.

### (i) (Overheads)



### (ii) (Rebalancing Overheads)

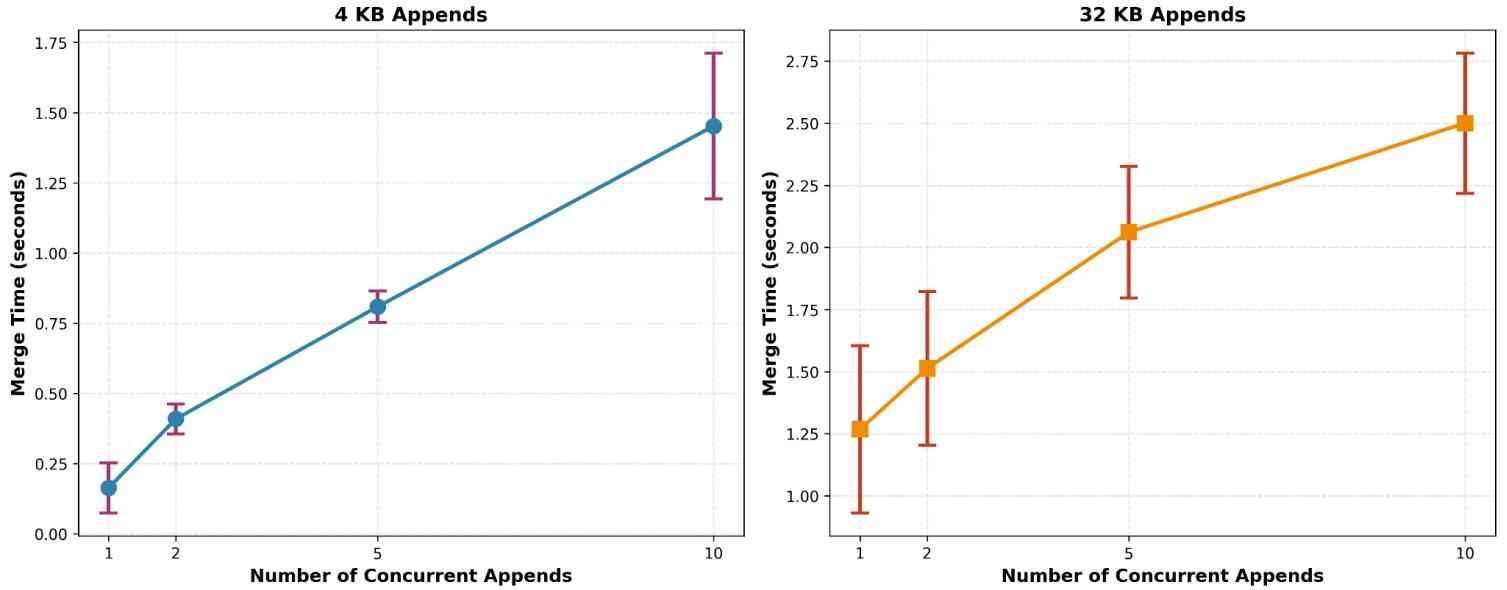


# CS425 Distributed Systems MP3 Report

Aditi Sadwelkar, Aryaman Ramchandran

## (iii) (Merge Performance)

Merge Performance by Append Size



## (iv) (Subsequent-merge Performance)

Merge Performance Comparison: First vs Second Merge  
(Second merge ~27% faster due to already-sorted data)

